

@workspace I am starting a new project called "GLP-1 Regulatory Intelligence Platform". I need you to act as the Lead Architect and Senior Python Developer. I will provide the full system design below. Your job is to understand the entire architecture, data flow, and constraints, and then help me implement it file-by-file starting with the data models.

## 1. Project Overview

We are building a Regulatory Intelligence Dashboard for analyzing and comparing the FDA labels of 10 specific GLP-1 drugs (e.g., Ozempic, Mounjaro, Wegovy).

- **Core Value:** Users can view parsed label text, see AI-extracted entities (Dosage, Side Effects), and use a Semantic Search Chatbot (RAG) to ask questions like "Does Ozempic cause thyroid tumors?".
- **Key Constraint:** Reliability and Deployment-Readiness. We use an "Offline-First" ingestion strategy.

## 2. System Architecture (The "Bakery" Model)

The system is divided into two phases: **Phase A (Ingestion)** and **Phase B (Application)**.

### *Phase A: The Data Pipeline (Offline/Background)*

- **Source:** We manually download 10 raw XML .zip files from FDA DailyMed and upload them to an AWS S3 Bucket (Cold Storage).
- **Ingestion Logic (builder.py):**
  - **Fetch:** Script pulls the raw zip from S3.
  - **Parse:** Uses lxml and specific LOINC codes (e.g., 34067-9 for Indications) to extract clean, verbatim text.
  - **Enrich (NER):** Uses a BioBERT model to extract structured entities (Strength, Route, Frequency) into JSON.
  - **Vectorize:** Chunks the text and generates embeddings using SentenceTransformer for semantic search.
- **Storage (Polyglot Persistence):**
  - **PostgreSQL:** Stores the clean text, metadata, and JSONB NER data. (Primary Key: SetID + Version).
  - **Vector DB (Pinecone/Chroma):** Stores the embeddings with metadata pointing back to the Postgres SetID.

## **Phase B: The Application (Online/Runtime)**

- **Backend:** FastAPI. It is "Read-Heavy."
  - GET /drugs/{id}: Fetches text from Postgres.
  - POST /chat: Performs RAG (Vector Search -> LLM Synthesis).
- **Frontend:** React (Vite) + Tailwind.
  - **Workspace:** A Rich-Text Reader view (not PDF) with a Split-Pane layout for side-by-side comparison.
  - **Sticky Chat:** A floating RAG widget that is context-aware (filters vector search by the currently viewed Drug ID).

## **3. Implementation Rules (Strict)**

1. **Storage:** Use AWS S3 (boto3) for raw files. Use PostgreSQL (via Docker) for app data.
2. **Versioning:** Use an "Append-Only" strategy. If a drug updates, we insert a new row with Version: 2 and archive Version: 1.
3. **Comparison:** \* **Lexical (Red-Lining):** Compares the CLEAN TEXT from Postgres (not raw XML).
  - a. **Semantic:** Uses Vector similarity.
4. **Tech Stack:** Python 3.10+, FastAPI, Pydantic v2, SQLAlchemy (Async), React, Docker.

## **4. Immediate Execution Plan**

We are starting with **Phase 1: The Blueprint**. I need you to help me set up the project structure and strictly define the Data Models before we write any logic.

**Task 1:** Create the directory structure:

```
```text
glp1_project
  -- backend
    -- core
      -- config
      -- models (pydantic & sqlalchemy)
      -- services
        -- s3
        -- etl
        -- vector
    -- data (local temp)
  -- docker-compose.yml
```

```

**Task 2:** Create backend/models/schemas.py. We need Pydantic models to define the shape of our data.

- DrugSection: title (str), text (str), loinc\_code (str).
- DrugMetadata: set\_id (UUID), name, manufacturer, version (int), upload\_date.
- NEREntity: label (str), text (str), confidence (float).

- DrugLabel: The master object containing metadata + list of sections + list of entities.

**Task 3:** Create docker-compose.yml.

- Service 1: postgres:15-alpine with a named volume for persistence.
- Service 2: pgadmin (optional).

**Task 4:** Create backend/services/s3\_client.py.

- A robust class to handle downloading/uploading to the bucket defined in .env.

## References :

This is the **Execution Roadmap**. We are moving from "Planning" to "Doing."

Since we are following the "**Reliability-First**" approach, we will not write a single line of Python until the Infrastructure (S3 + Databases) is ready to receive data.

Here is your broad, step-by-step guide to starting the implementation.

### Phase 0: The Tooling Setup (Do this First)

Before handling data, you need the tools to hold it.

1. **AWS Account:** Sign up (Free Tier is fine). Create an **S3 Bucket** (e.g., glp1-project-raw-v1).
  - a. *Action:* Go to IAM (Identity & Access Management) and create a User with AmazonS3FullAccess.
  - b. *Save:* The ACCESS\_KEY and SECRET\_KEY. You will need these to let your Python script talk to S3.
2. **PostgreSQL Database:**
  - a. *Recommendation:* Since you want this "Deployment Ready," use **Docker** on your local machine. It mirrors how it will run in the cloud.

- b. *Alternative:* Use a managed free tier like **Supabase** or **Neon** if you don't want to manage Docker containers yet.

### 3. Vector Database:

- a. *Recommendation:* Sign up for **Pinecone** (easiest for beginners) or **ChromaDB** (runs locally). Pinecone is better for the "Deployment" goal as it persists online.

## Phase 1: The "Raw" Data Acquisition

**Goal:** Get the "Golden Set" of 10 labels safely into the Cloud Archive.

### Step-by-Step Actions:

#### 1. Manual Download (The Source):

- a. Go to **DailyMed** ([dailymed.nlm.nih.gov](http://dailymed.nlm.nih.gov)).
- b. Search for "Ozempic".
- c. Look for the button/tab: "**Download Mapping / SPL Resources**".
- d. Download the **XML (zip)** file. Do NOT unzip it.
- e. Repeat for all 10 drugs (Mounjaro, Wegovy, etc.).
- f. *Result:* You have 10 .zip files on your laptop.

#### 2. The S3 Injection (The Archive):

- a. *Method A (Easy):* Go to the **AWS Console** website → S3 → Your Bucket. Click "**Upload**" and drag-and-drop the 10 zip files.
- b. *Method B (Pro):* Configure the AWS CLI on your terminal (`aws configure`) and run `aws s3 cp ./local_folder s3://your-bucket-name --recursive`.
- c. *Result:* Your raw data is now "Cloud Native." It is safe.

## Phase 2: The "Builder" Environment

**Goal:** Set up the Python logic that will read S3 and write to Postgres.

You are not building a "Web App" yet. You are building a **Data Pipeline Script**.

1. **The Project Folder:** Create a folder named `glp1_intelligence`. Inside, create this structure:

```
Plaintext
/glp1_intelligence
  /backend
    /core
      config.py    # Store AWS Keys / DB URLs here
    /etl        # "Extract Transform Load"
      builder.py   # <--- THE ENGINE. This is your main script.
      parser.py    # The lxml logic
      ner.py       # The AI logic
      schemas.py   # The Pydantic blueprints
  /frontend     # Leave empty for now
```

**2. The Pipeline Logic (How it flows):** You will write the builder.py script to do the following (conceptually):

- a. **Connect:** "Hello S3, give me the list of files."
- b. **Download to Memory:** Pull ozempic.zip into RAM (don't save to disk).
- c. **Shred:** Pass the file to parser.py. It returns clean text.
- d. **Analyze:** Pass the text to ner.py. It returns JSON entities + Vectors.
- e. **Push:**
  - i. "Hello Postgres, insert this text."
  - ii. "Hello Pinecone, insert these vectors."

## Phase 3: The Database Schema (The Destination)

**Goal:** Prepare Postgres to receive the data.

You need to create the Tables *before* running the script.

### 1. The drug\_labels Table:

#### a. Columns:

- i. set\_id (Primary Key, String/UUID) - *The Link*.
- ii. version (Integer) - *For history*.
- iii. name (String) - e.g., "Ozempic".
- iv. manufacturer (String).
- v. content\_json (JSONB) - *Stores the sections: { "indications": "...", "dosage": "..." }*.

- vi. ner\_data (JSONB) - Stores entities: { "strength": "5mg" }.
- vii. is\_active (Boolean) - True for latest version.

## 1. Single Label Analysis (The Workspace View)

When the user clicks "Ozempic," they do **not** see a wall of text.

- **The Layout:** We use a **Split-Pane Layout**.
  - **Left Sidebar (Navigation):** A vertical list of sections derived from the LOINC codes (e.g., 1. *Indications*, 2. *Dosage*, 3. *Warnings*, 4. *Adverse Reactions*).
  - **Main Content Area (The Reader):** This displays the **Cleaned Text** (from PostgreSQL).
- **The Experience:**
  - When the user clicks "Dosage" on the left, the Main Content automatically scrolls to that specific section.
  - **NER Highlights:** Inside this text, specific entities (like 0.5 mg) are highlighted with small colored "chips" or background colors to make them pop.
- **Data Source:** This view is purely rendering the **JSONB** content stored in PostgreSQL. It is *not* reading XML.

## 2. The Comparison View (The "Diff" Tool)

This is where your "Red Lining" and "Color Coding" come into play. The user selects two drugs (e.g., Ozempic vs. Mounjaro) and enters **Comparison Mode**.

We display this as **Two Columns Side-by-Side**.

### A. *Lexical Comparison (The "Red Lining")*

- **What is it?** A standard "Word-for-Word" check (like Microsoft Word "Track Changes").
- **Where does it happen?** On the **Cleaned Text Strings** from PostgreSQL.
- **The Visuals:**
  - **Red Strike-through:** Text that exists in Drug A but *not* in Drug B.

- **Green Underline:** Text that exists in Drug B but *not* in Drug A.
- **Black:** Text that is identical.
- **Use Case:** Best for comparing **Version 1 vs. Version 2** of the *same* drug (to see exactly what words the FDA changed).

### **B. Semantic Comparison (The "Traffic Light" Segregation)**

- **What is it?** A "Meaning" check. This uses the **Vectors** to calculate similarity scores between paragraphs.
- **Where does it happen?** It overlays color blocks on top of the **Cleaned Text Paragraphs**.
- **The Visuals:**
  -  **Green Background:** "High Similarity." (e.g., Ozempic says "Nausea"; Mounjaro says "Nausea". The words match or are synonyms).
  -  **Yellow Background:** "Partial Match." (e.g., Ozempic says "Thyroid Tumors in Rats"; Mounjaro says "Thyroid Tumors in Mice". The concept is similar but details differ).
  -  **Red Background:** "Unique / Conflict." (e.g., Ozempic says "Inject Weekly"; Rybelsus says "Take Daily". These are fundamentally different concepts).
- **Use Case:** Best for comparing **Drug A vs. Drug B** (Competitor Analysis).

### **The Golden Rule of this Architecture**

**AWS S3 is for "Cold Storage" (The Archive). Databases (Postgres/Vector) are for "Live Access" (The App). The Frontend NEVER talks to S3 directly.**

### **1. AWS S3: The "Raw Archive"**

- **What do we store here?**
  - We store the **Raw Zipped XML files** exactly as they came from the FDA.
- **Why?**

- It is our "Source of Truth." If we mess up our parsing logic later, we can delete our database and re-process everything from S3 without needing to download from the FDA again.
- **Format:** Raw .zip or .xml. **NOT JSON.**
- **Access Pattern:**
  - **Who accesses it?** Only the **Backend Python Script** (builder.py).
  - **When?** Only during the **Ingestion Phase** (Offline).
  - **Does the User see this?** NO. The dashboard never reads from S3.

## 2. The Processing Logic (The "Bridge")

After the file is in S3, your builder.py script wakes up. It reads the XML from S3 into memory. It does **three things** to that data:

1. **Parses Text:** Extracts clear text (Strings) from the XML tree.
2. **Extracts Entities (NER):** Finds "5mg", "Nausea", etc.
3. **Vectorizes (Embeddings):** Converts text chunks into Numbers ([0.1, 0.5...]).

Now, it sends this processed data to **Two Different Databases**.

## 3. Database A: PostgreSQL (The "Structured Store")

- **What is stored here?**
  - **Metadata:** Drug Name, SetID (Primary Key), Manufacturer.
  - **The Clean Text:** The actual paragraphs of the label (Indications, Dosage) as long strings.
  - **NER Output:** The structured JSON output from your NER model (e.g., {"strength": "5mg", "side\_effects": ["Nausea"]}).
- **Why?**
  - This is what populates the **Dashboard** and the **Workspace**.
- **Access Pattern:**
  - **Who accesses it?** The FastAPI Backend (GET /drugs/{id}).
  - **Scenario:** When the user clicks "Ozempic," the API queries **Postgres** to get the text to display on the screen. **We do NOT display raw XML from S3.** We display the clean text from Postgres.

## 4. Database B: Vector DB (The "Semantic Brain")

- **What is stored here?**
  - The **Vectors** (Embeddings) derived from the text chunks.
  - **Metadata:** drug\_id (Foreign Key) and the text\_content of that chunk.
- **Why?**
  - This is EXCLUSIVELY for the **Search Bar** and **RAG Chatbot**.
- **Access Pattern:**
  - **Scenario:** User searches "drugs for stomach pain."
  - The system queries the Vector DB to find mathematically similar concepts.

## 5. The "Link": Primary Key vs. Foreign Key

You asked how we ensure accuracy. We use the **SetID** (The FDA's unique ID for the drug).

- **In PostgreSQL:** SetID is the **Primary Key**.
  - Row: ID: 123 | Name: Ozempic | Text: "..."
- **In Vector DB:** SetID is stored as **Metadata (Foreign Key)** attached to every single vector.
  - Vector: [0.1, 0.9...] | Metadata: { "drug\_id": "123", "chunk\_id": 1 }

### How RAG uses this:

1. User acts on **Ozempic** (ID: 123).
2. User asks a question.
3. Backend searches Vector DB: *"Find vectors close to this question, BUT... FILTER WHERE drug\_id == 123."*
4. This guarantees the AI never reads Mounjaro's data while you are looking at Ozempic.

## Detailed Scenario Walkthrough

Let's trace the exact flow so you are 100% confident.

### **Step 1: The User Opens the Dashboard**

- **User sees:** A list of 10 drugs (Ozempic, Mounjaro, etc.).
- **Backend Action:** SELECT name, manufacturer, set\_id FROM drugs\_table; (PostgreSQL).
- **S3 Usage:** None.
- **Vector DB Usage:** None.

### **Step 2: The User Clicks "Ozempic" (The Workspace)**

- **User sees:** The full label text, divided into tabs (Indications, Dosage). They see "0.5mg" highlighted (NER data).
- **Backend Action:** SELECT content, ner\_data FROM drugs\_table WHERE set\_id = '123'; (PostgreSQL).
- **Clarification:** We display the **Clean Text** from Postgres. We do *not* fetch the raw XML from S3. The XML is ugly and hard to read. The Postgres text is already cleaned/stripped of HTML tags during the processing phase.

### **Step 3: The User Asks a Question (The Chat)**

- **User types:** "Does it cause cancer?"
- **Backend Action:**
  - Convert question to Vector.
  - Query **Vector DB:** Find chunks matching "cancer" where drug\_id = '123'.
  - Send those chunks to LLM.
  - Return answer.

## **1. The Revised Storage Strategy (Cloud-Native)**

Instead of a data/raw folder on your laptop, we will use **AWS S3 Buckets**.

- **Bucket Name:** glp1-raw-labels
- **Structure:**

Plaintext

```
s3://glp1-raw-labels/  
/ozempic/  
v1_2024.zip  
v2_2025.zip <-- (Current)  
/mounjaro/  
v3_2025.zip
```

### Why this is professional:

1. **Accessible Anywhere:** Your deployed API (running on a server) and your local Dev environment can both access the exact same files.
2. **Versioning:** S3 has built-in versioning. You never accidentally lose an old label.

## 2. The Automated "Watchdog" Pipeline (Deep Dive)

You asked: *How do we replace older versions without user interaction? Is it complicated?*

**The Answer:** It is **not complicated** if we follow a strict "Check-Compare-Swap" logic. We write a single script (watchdog.py) that runs automatically (e.g., once every 24 hours).

Here is the exact logic flow:

### Step A: The Poll (The "Check")

The script wakes up. It holds a list of the 10 drugs we care about (specifically their **SetIDs**).

- **Action:** It loops through the list and hits the FDA DailyMed API history endpoint.
  - *Request:* GET /spl/{set\_id}/history
  - *Response:* "The latest version for Ozempic is **Version 5**, published **Jan 20, 2026**."

### Step B: The Comparison (The "Detect")

The script looks at your **Postgres Database**.

- *Query:* SELECT version FROM drugs WHERE name = 'Ozempic'
- *Result:* "We currently have **Version 4**."

- **Logic:** Version 5 > Version 4. **UPDATE REQUIRED.**

### ***Step C: The "Silent Swap" (The Execution)***

This happens entirely in the background. The user sees nothing until it is finished.

1. **Download:** The script downloads the new Version 5 zip file from the FDA.
2. **Upload to S3:** It saves it to s3://glp1-raw-labels/ozempic/v5.zip.
3. **Process (The Oven):** The script triggers the parsing logic (just like we did manually on Day 1).
  - a. Parses XML -> Extracts NER -> Creates Vectors.
4. **Atomic Update (The Switch):** This is the crucial part. We update the database records.
  - a. *Postgres:* Update the metadata rows with the new text.
  - b. *Vector DB:* **Delete** old vectors for Ozempic and **Insert** the new ones.
  - c. *Status:* Update the "Last Updated" timestamp.

**Is it complicated?** No. It is just a function that runs in a loop.

- **User Interaction:** None.
- **Downtime:** None. The old data remains available until the millisecond the new transaction commits.

## **3. The Revised Implementation Path**

Since we are now using S3, your "Starting Point" shifts slightly. You need to set up the plumbing first.

### ***Step 1: The Setup (Infrastructure)***

- **AWS:** Create an S3 bucket (e.g., my-pharma-project-raw). Get your AWS\_ACCESS\_KEY and AWS\_SECRET\_KEY.
- **Local Env:** Configure your project to talk to S3 using the boto3 library (standard Python SDK for AWS).

### **Step 2: Manual Ingestion (Upload to Cloud)**

- **Action:** Download the 10 zip files manually (as discussed).
- **Change:** Instead of saving them to a local folder, write a small script (upload\_initial.py) to upload them to your S3 bucket.
- **Why?** This mimics the production environment from Day 1.

### **Step 3: The Blueprint (schemas.py)**

- (Same as before) Define your data models.

### **Step 4: The Processor (builder.py)**

- **Change:** Modify the script so it reads **from S3**, not from your hard drive.
  - **Code:** s3\_client.download\_file(...) → Parse → Push to DB

## **The User Journey (The "Story")**

**The User:** Jane, a Regulatory Intelligence Analyst. **The Goal:** Compare the safety profiles of **Ozempic** (Novo Nordisk) and **Mounjaro** (Eli Lilly).

### **Step 1: The Global Dashboard (Instant Access)**

- Jane logs in. She sees a clean table listing the 10 monitored GLP-1 drugs.
- **Visuals:** Each row shows the Drug Name, Manufacturer, and a green badge: “*Latest Update: Jan 15, 2026 (Version 3)*”.
- **Action:** She clicks on “**Ozempic**”.

### **Step 2: The Workspace (Deep Analysis)**

- The screen transitions to the **Single Drug View**.
- **Layout:**
  - **Left Sidebar:** Navigation menu (1. Indications, 2. Dosage, 3. Warnings...).

- **Center:** The cleaned, easy-to-read label text (served from PostgreSQL). It looks like a modern document, not a PDF.
- **Highlights:** She sees small blue chips floating over text like "*0.25 mg*" or "*Nausea*". These are the **NER Entities** our AI extracted automatically.
- **Action:** She wants to check for thyroid risks, but doesn't want to read 50 pages.

### **Step 3: The Sticky Chat (Semantic Search)**

- She clicks the **Floating Chat Icon** (bottom-right).
- **Query:** "*Does this drug have a boxed warning for thyroid tumors?*"
- **The System (RAG):**
  - Searches the **Vector Database** specifically for Ozempic chunks.
  - Finds the "Boxed Warning" section.
  - **AI Response:** "*Yes. Ozempic causes thyroid C-cell tumors in rodents. It is contraindicated in patients with a family history of MTC.*"
  - **Citation:** The AI provides a clickable link. Jane clicks it, and the center screen auto-scrolls to the exact paragraph in the "Warnings" section.

### **Step 4: The Comparison (Side-by-Side)**

- Jane returns to the Dashboard and selects "**Compare**". She picks **Ozempic** (Left) and **Mounjaro** (Right).
- **Lexical View (Red-Lining):** She sees the text of both. Words present in Ozempic but missing in Mounjaro are highlighted in **Red**; new words in Mounjaro are **Green**.
- **Semantic View:** She switches to "Concept Mode."
  - The "Adverse Reactions" sections for both align perfectly side-by-side.
  - The system highlights "Nausea" in Green on both sides (indicating a match), even though Ozempic writes "nausea" and Mounjaro writes "nausea and vomiting"

## **The Tech Stack & Component Guide**

This is the inventory of tools you need to install.

### **1. The Infrastructure (The Foundation)**

- **AWS S3 (Boto3): The Raw Archive.**
  - **Usage:** Stores the original .zip files downloaded from FDA.

- *Why:* Safety. If we break the database, we rebuild from S3.
- **Docker: The Container Engine.**
  - *Usage:* Runs the PostgreSQL database cleanly on your laptop without "installing" it in Windows/Mac.
  - *Why:* Deployment readiness.

## **2. The Databases (Polyglot Persistence)**

- **PostgreSQL (v15+): The Structured Store.**
  - *Usage:* Stores the Clean Text (for the dashboard), Metadata (SetIDs), and NER JSON.
  - *Why:* We need ACID compliance and reliability for the main app data.
- **Pinecone (or ChromaDB): The Vector Store.**
  - *Usage:* Stores the Embeddings (Numbers) for the Chatbot.
  - *Why:* Standard SQL databases cannot perform "Semantic Search" (finding similar meanings).

## **3. The Processing Core (Python)**

- **lxml: The Parser.**
  - *Usage:* Cuts the XML tree into sections using LOINC codes.
  - *Why:* Faster and more accurate than standard Regex or BeautifulSoup for XML.
- **HuggingFace Transformers (BioBERT): The Extractor.**
  - *Usage:* Named Entity Recognition (NER) to find specific medical variables.
- **SentenceTransformers: The Embedder.**
  - *Usage:* Converts text to vectors (model.encode(text)).

## **4. The Application Layer**

- **FastAPI: The Backend API.**
  - *Usage:* Connects the Frontend to the Databases. Handles the "Chat" streaming.
- **React + Tailwind CSS: The Frontend UI.**
  - *Usage:* The Dashboard, Split-Pane view, and Chat Widget.

## The Implementation Guide (Step-by-Step)

Follow this sequence strictly to avoid backtracking.

### **Phase 1: Setup & Infrastructure (Day 1)**

1. **Project Init:** Create the folder structure (backend/, frontend/, data/).
2. **AWS S3:** Create a bucket. Manually upload your 10 FDA Zip files.
3. **Docker:** Create docker-compose.yml.

YAML

```
services:  
  db:  
    image: postgres:15-alpine  
    volumes:  
      - ./postgres_data:/var/lib/postgresql/data  
    ports:  
      - "5432:5432"
```

*Run docker-compose up -d. You now have a database.*

### **Phase 2: Data Modeling (Day 2)**

- **Goal:** Define the "Contract."
- Create backend/schemas.py.
  - Define DrugMetadata (name, set\_id, version).
  - Define DrugSection (title, content\_text).
  - Define DrugNER (strength, route).

### **Phase 3: The ETL Pipeline (Day 3-4)**

- **Goal:** Move data from S3 → Postgres/Vector DB.
- Create backend/builder.py.
  - **Fetch:** Connect to S3 (boto3) and download ozempic.zip to memory.

- **Parse:** Use lxml to extract text. *Validation:* Print the text to console to ensure it's clean.
- **Vectorize:** Generate embeddings for the text.
- **Load:** Connect to Postgres (using SQLAlchemy) and Pinecone. Insert the data.
- *Result:* Your database is now populated with real data.

#### ***Phase 4: The Backend API (Day 5)***

- **Goal:** Expose the data.
- Create backend/main.py.
  - GET /drugs: Select \* from Postgres.
  - GET /drugs/{id}: Return the full text for the Workspace.
  - POST /chat: Receive query → Pinecone Search → LLM API → Return Answer.

#### ***Phase 5: The Frontend (Day 6-7)***

- **Goal:** Visuals.
- Initialize React (npm create vite@latest).
- **Dashboard:** Fetch list from API.
- **Workspace:** Render the HTML text. Use divs for sections.
- **Chat:** Add the floating button component.

#### ***Phase 6: The Watchdog (Bonus/Final)***

- Create watchdog.py.
- Script checks FDA API for new versions. If found, it triggers builder.py to append the new version to the DB.